

Q1: Gaussian Mixture Model(GMM)

(a) GMM learning with Expectation Maximization

GMM with EM on 1D data

```
In [1]: import matplotlib.pyplot as plt
from sklearn import cluster, datasets, mixture
import numpy as np
from scipy.stats import multivariate_normal
```

Creating the 1d dataset

```
In [63]: # define the number of points
n_samples = 100
mu1, sigma1 = -4, 1.2 # mean and variance
mu2, sigma2 = 4, 2.2 # mean and variance
mu3, sigma3 = 0, 1.6 # mean and variance

x1 = np.random.normal(mu1, np.sqrt(sigma1), n_samples)
x2 = np.random.normal(mu2, np.sqrt(sigma2), n_samples)
x3 = np.random.normal(mu3, np.sqrt(sigma3), n_samples)

X = np.array(list(x1) + list(x2) + list(x3))
np.random.shuffle(X)
print("Dataset shape:", X.shape)
```

Dataset shape: (300,)

```
In [64]: def pdf(data, mean: float, variance: float):
# A normal continuous random variable.
# Enter your code here 1

return np.exp(-(data-mean)**2/(2*variance))/(np.sqrt(2*np.pi*variance)+eps)
```


```
In [65]: # visualize the training data
bins = np.linspace(np.min(X), np.max(X), 100)

plt.figure(figsize=(10,7))
plt.xlabel("$x$")
plt.ylabel("pdf")
plt.scatter(X, [0.005] * len(X), color='navy', s=30, marker=2, label="Train data")

plt.plot(bins, pdf(bins, mu1, sigma1), color='red', label="True pdf")
plt.plot(bins, pdf(bins, mu2, sigma2), color='red')
plt.plot(bins, pdf(bins, mu3, sigma3), color='red')

plt.legend()
plt.plot()
```

Out[65]: []



```
In [90]: # define the number of clusters to be learned
k = 3
weights = np.ones((k)) / k
means = np.random.choice(X, k)
variances = np.random.random_sample(size=k)
print(means, variances)

[-5.16860852  0.13606389  2.94698077] [0.99164432 0.69381628 0.0612628 ]
```

Actual implementation of EM

```

In [91]: eps=1e-8
for step in range(10):

    if step % 1 == 0:
        plt.figure(figsize=(10,6))
        axes = plt.gca()
        plt.xlabel("$x$")
        plt.ylabel("pdf")
        plt.title("Iteration {}".format(step))
        plt.scatter(X, [0.005] * len(X), color='navy', s=30, marker='x', label="Train data")

        plt.plot(bins, pdf(bins, mu1, sigma1), color='grey', label="True pdf")
        plt.plot(bins, pdf(bins, mu2, sigma2), color='grey')
        plt.plot(bins, pdf(bins, mu3, sigma3), color='grey')

        plt.plot(bins, pdf(bins, means[0], variances[0]), color='blue', label="Cluster 1")
        plt.plot(bins, pdf(bins, means[1], variances[1]), color='green', label="Cluster 2")
        plt.plot(bins, pdf(bins, means[2], variances[2]), color='magenta', label="Cluster 3")

        plt.legend(loc='upper left')

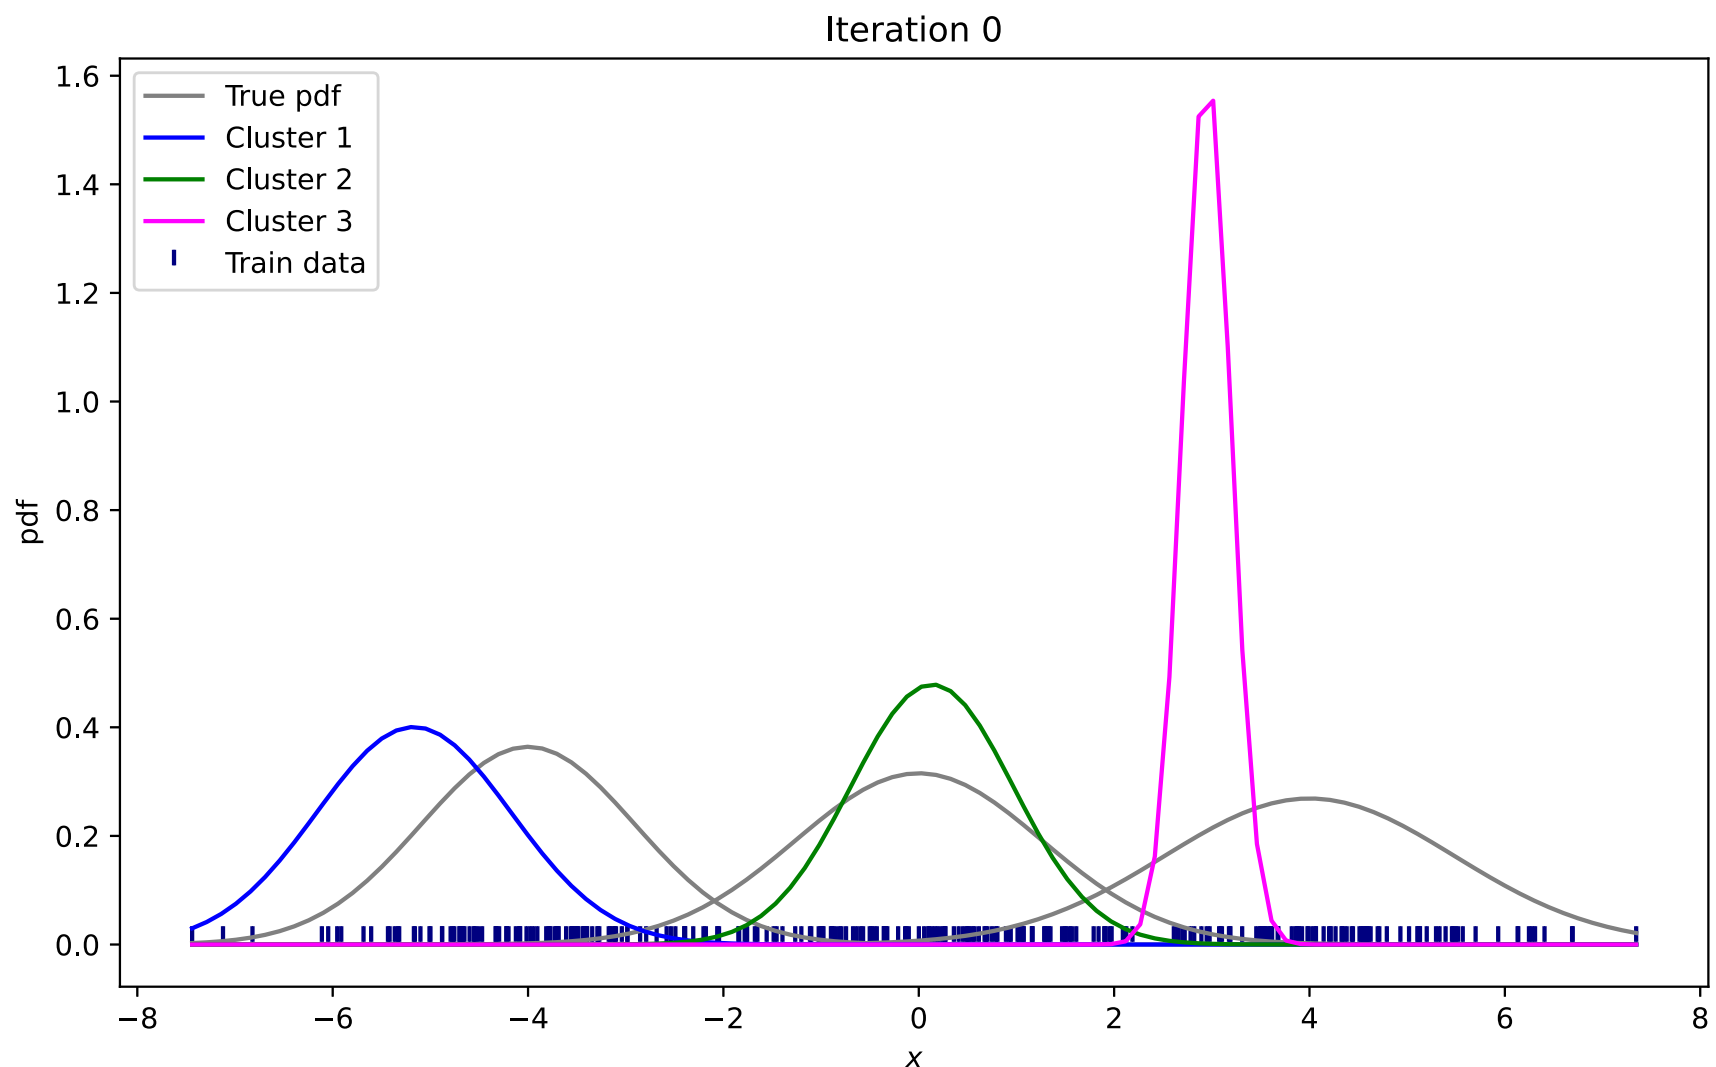
        plt.savefig("img_{0:02d}".format(step), bbox_inches='tight')
        plt.show()

    # calculate the maximum likelihood of each observation xi
    likelihood = np.empty((300,0))

    # Expectation step
    for j in range(k):
        Pdf = pdf(X.reshape(-1,1), means[j], np.sqrt(variances[j]))
        likelihood = np.hstack((likelihood, Pdf))
    # likelihood = np.array(likelihood)
    print('likelihood', likelihood.shape)

    # b = []
    # Maximization step
    # Enter your code here 2
    a = likelihood*weights
    Delimiter = np.sum(a,axis = 1).reshape(-1,1)
    b = a/Delimiter
    means = (X.reshape(1,-1)@b/np.sum(b,axis = 0))[0]
    variances = np.sum(((X.reshape(-1,1)- means)**2)*b,axis = 0)/np.sum(b,axis = 0)
    weights = np.sum(b,axis = 0)/300
    print("means: ",means)
    print("variances: ",variances)
    print("weights: ",weights)


```




```

likelihood (300, 3)
means: [-4.17187097  0.13129718  4.0200934 ]
variances: [1.08147581  3.19699756  1.2478432 ]
weights: [0.31893197  0.36723472  0.31383331]


```




```
likelihood (300, 3)
means: [-4.10872796  0.03154118  4.25508907]
variances: [1.21739936 1.82147395 1.42109413]
weights: [0.32517156 0.36843446 0.30639398]
```




```
likelihood (300, 3)
means: [-4.06564106  0.0566804  4.25328061]
variances: [1.26579359 1.56647465 1.38497272]
weights: [0.3341808 0.35598904 0.30983016]
```




```
likelihood (300, 3)
means: [-4.03418774  0.08217868  4.2453569 ]
variances: [1.30773111 1.46961666 1.3909682 ]
weights: [0.34000701 0.34851753 0.31147545]
```




```
likelihood (300, 3)
means: [-4.01298117  0.10060161  4.24084261]
variances: [1.33854001 1.42083011 1.39713468]
weights: [0.34374127 0.34402579 0.31223294]
```



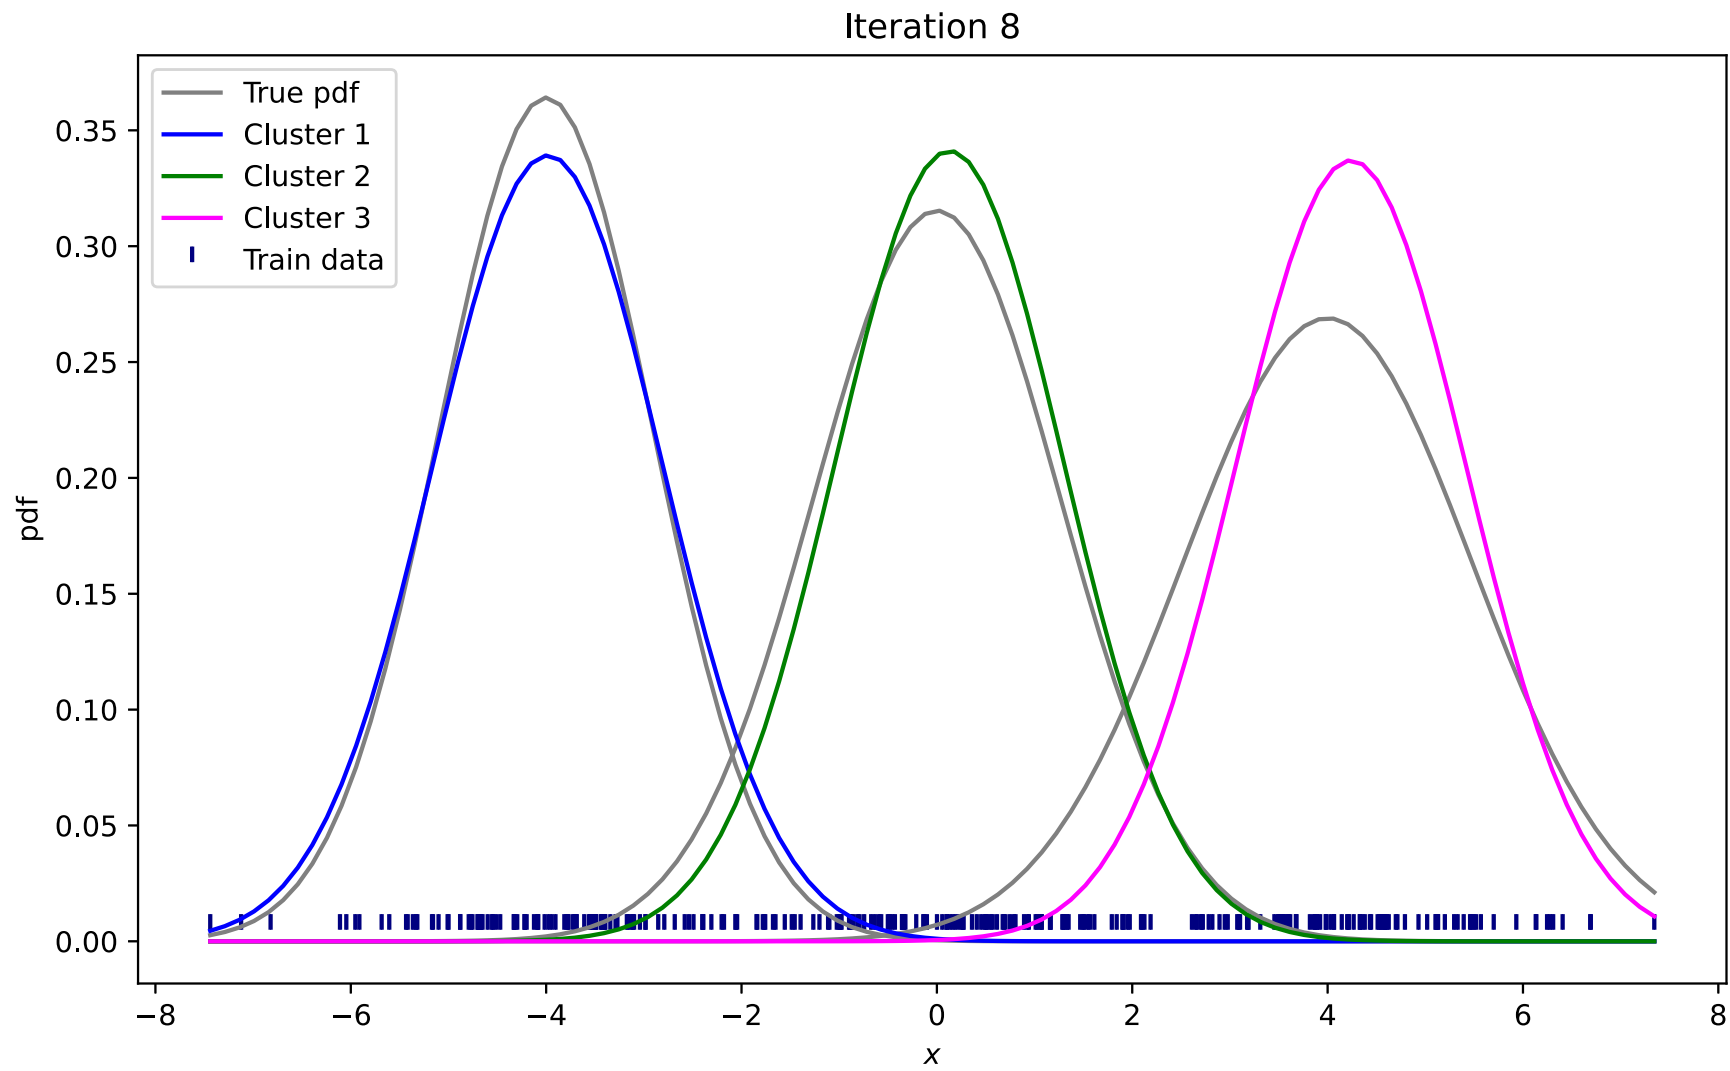
```
likelihood (300, 3)
means: [-3.99901235  0.11343548  4.23890387]
variances: [1.3596408  1.39326567  1.40002673]
weights: [0.34614828  0.34130677  0.31254495]
```



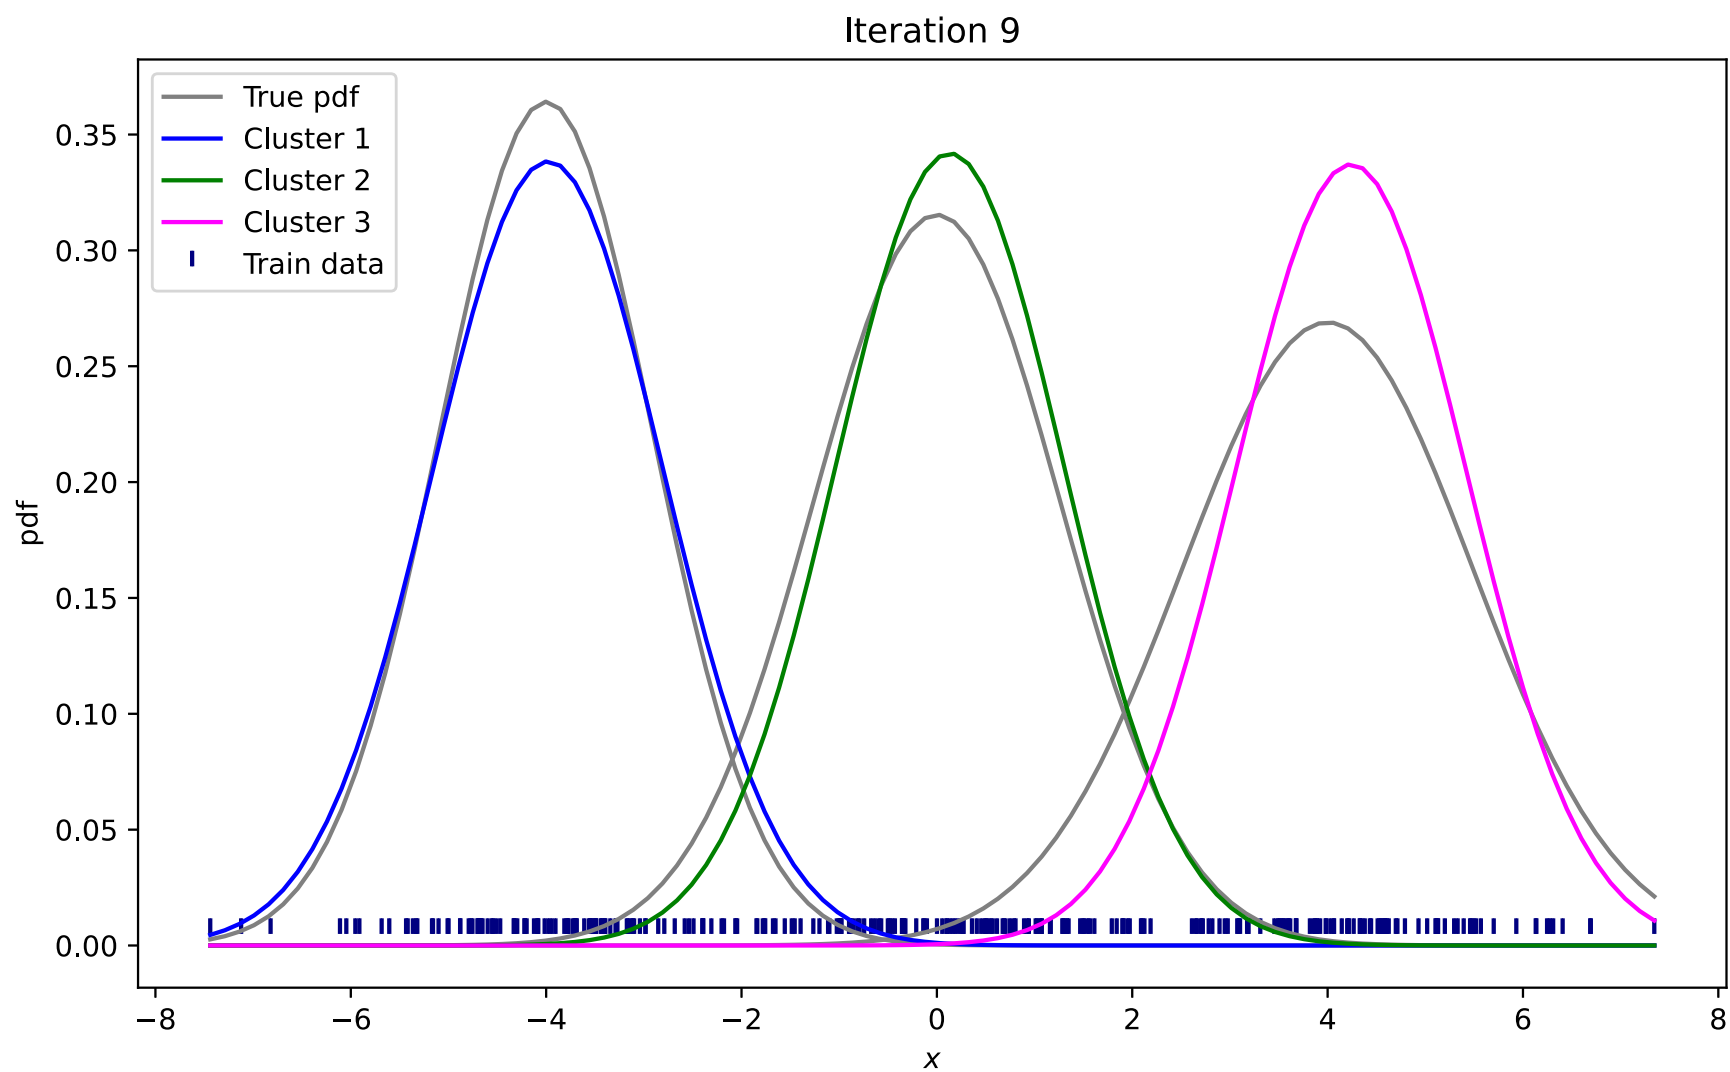
```
likelihood (300, 3)
means: [-3.98985141  0.1223494  4.23834836]
variances: [1.3737631  1.37700041  1.40084175]
weights: [0.34770982  0.33965252  0.31263765]
```



```
likelihood (300, 3)
means: [-3.98383508  0.12854391  4.23844037]
variances: [1.38314827  1.36721716  1.4006491 ]
weights: [0.348729  0.33864147  0.31262952]
```



```
likelihood (300, 3)
means: [-3.97987145  0.13284944  4.23878688]
variances: [1.38937716  1.36126811  1.40006091]
weights: [0.34939785  0.3380208  0.31258135]
```



```
likelihood (300, 3)
means: [-3.97725066  0.13584095  4.23919365]
variances: [1.39351551 1.35762205 1.39938458]
weights: [0.34983897 0.33763809 0.31252294]
```

(b) GMM with sklearn library

```
In [92]: from sklearn.mixture import GaussianMixture
gm = GaussianMixture(n_components=3, random_state=0).fit(X.reshape(-1,1))
means = gm.means_
variances = gm.covariances_
print("means: ",means)
print("variances: ",variances)

means: [[ 4.24667335]
 [-3.95912856]
 [ 0.16076671]]
variances: [[[1.43473826]]

 [[1.46749592]]


 [[1.55910445]]]
```

I get similar values compared to a. If we have more iterations in (a), we can get more robust values.

Q2: Clustering Algorithms


I. GMM

```
In [103... X_ = np.loadtxt("threeblobs.txt")
gm = GaussianMixture(n_components=3, random_state=0).fit(X_)
labels = gm.predict(X_)
# Plotting
plt.figure(figsize=(10, 8))
plt.scatter(X_[ :,0],X_[ :,1],c = labels)
plt.xlabel("x")
plt.ylabel("y")
plt.show()
```

II. DBSCAN

```
In [118... from sklearn.cluster import DBSCAN
clustering = DBSCAN(eps=0.6, min_samples=10).fit(X_)
labels = clustering.labels_
# Plotting
plt.figure(figsize=(10, 8))
plt.scatter(X_[ :,0],X_[ :,1],c = labels)
plt.xlabel("x")
plt.ylabel("y")
plt.show()
```



III. KMeans

```
In [119... from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=0).fit(X_)
labels = kmeans.labels_
# Plotting
plt.figure(figsize=(10, 8))
plt.scatter(X[:,0],X[:,1],c = labels)
plt.xlabel("x")
plt.ylabel("y")
plt.show()
```

